

Introduction to R

Alternative to our doc: [An Introduction to R by CRAN](#)

What is R and why R?

R (programming language)

From Wikipedia, the free encyclopedia

R is a [programming language](#) and [free software](#) environment for [statistical computing](#) and graphics supported by the R Core Team and the R Foundation for Statistical Computing.^[7] The R language is widely used among [statisticians](#) and [data miners](#) for developing [statistical software](#) and [data analysis](#). Polls, [data mining surveys](#), and studies of scholarly literature databases show substantial increases in R's popularity,^[8] since August 2021, R ranks 14th in the [TIOBE index](#), a measure of popularity of programming languages.^[9]

As you will see in the course, there are a lot of calculations and computations that are too difficult, too annoying, or outright impossible to do by hand.

Since the Year of Our Lord 2021, we can delegate the boring stuff to computers.

Q: How familiar/comfortable are you with programming?

Fx:

The sum of the primes below 10 is $2 + 3 + 5 + 7 = 17$.
Find the sum of all the primes below two million.

1. I know *exactly* how to do this
2. I know how to start but I can figure it out
3. I'm not sure but give me some time and I'll manage
4. I have no clue

Installation:

Follow installation instructions on [the R website](#). We will also use the [RStudio IDE](#).
Some minutes for troubleshooting.

Getting started

There is a command prompt

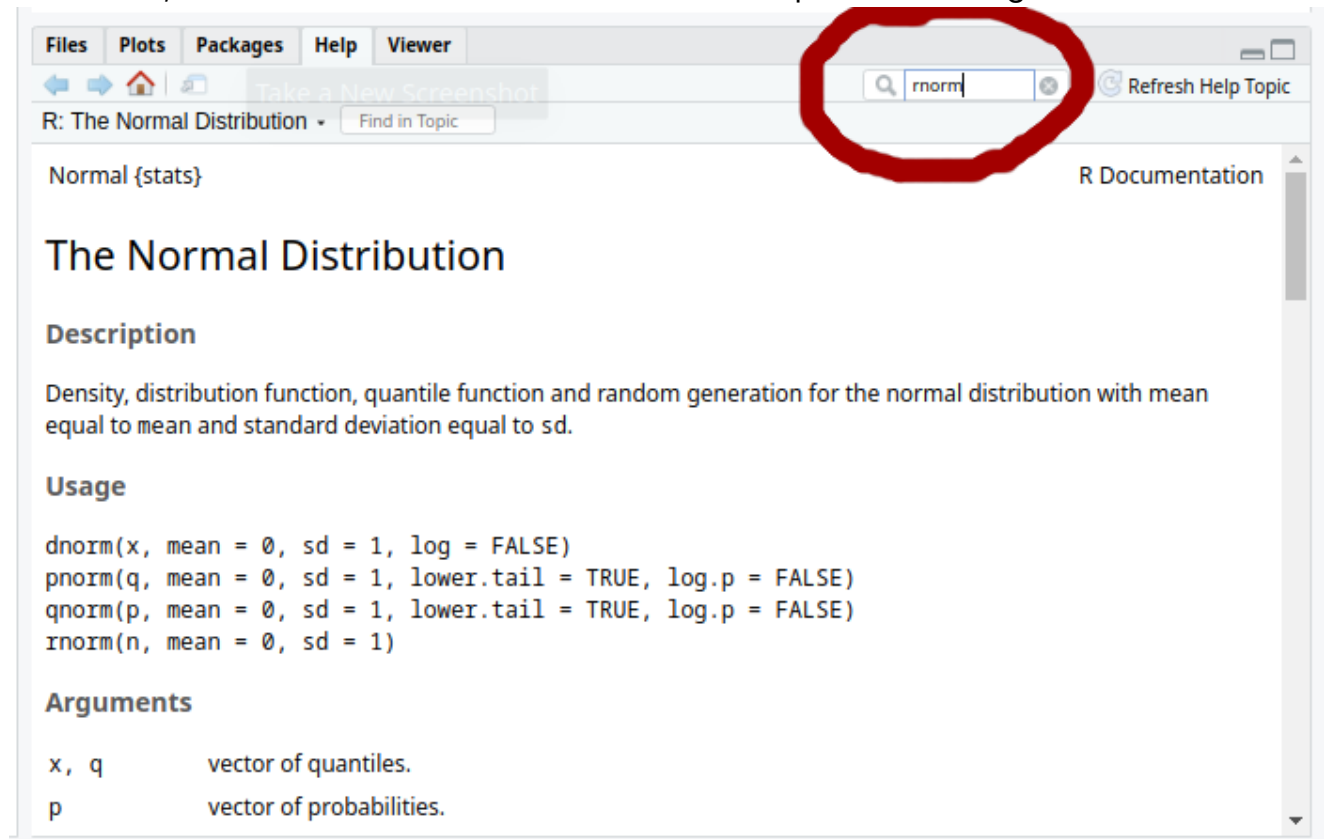
```
>
```

Here we type commands. We can ask for help:

```
help(anova)
```

If using RStudio, this opens the docpage in the window to the right.

In RStudio, we can also use the searchbar in the Help tab on the right:



We will talk about how to read this documentation below.

Basic syntax

Statements end with `;`, comments can be added with `#`.

Addition, multiplication, subtraction, division as per usual. Raising to some power is done with `a^b` or `a**b`, modulo function is `%%`.

Assignment of variables with `=`, `->` or `<-`:

```
> x <- 2
> 4 -> y
> x+y
[1] 6
```

[Why?](#)

Who hurt you, R

Note: The document states that we can use `_` to assign variables but I tried and didn't work (both RStudio and command-line). This was probably deprecated sometime between now and...

February 26, 2003

Oh.

Indeed, there is no mention of assignment of variables with `_` in the official Intro to R tutorial from CRAN.

Boolean values are `FALSE` and `TRUE` (also `0` for true and any non-zero for false).

You can list all the variables you have stored with `ls()`:

```
> ls()
[1] "x" "y"
```

And you can delete them with `rm`.

Arrays (vectors)

You can construct vectors with the `c` function:

```
> vec1 <- c(1,2,3,4)
> vec1
[1] 1 2 3 4
```

You can create arrays of consecutive integers with `min:max`:

```
> 5:10
[1] 5 6 7 8 9 10
```

Note that both limits are *inclusive*.

We can also create sequences with `seq`. The default syntax is

```
seq(min, max, step)
```

But we can pass the keyword argument `length` to set the length of the array instead of the step:

```
> seq(5,10,0.5) #same as seq(5,10,by=0.5)
[1] 5.0 5.5 6.0 6.5 7.0 7.5 8.0 8.5 9.0 9.5 10.0
> seq(5,10,length=8)
[1] 5.000000 5.714286 6.428571 7.142857 7.857143 8.571429 9.285714
[8] 10.000000
```

Note that the limits are *inclusive*.

Matrices:

Constructed from the unrolled version + number of rows or columns:

```
> matEntries <- 1:16
> matrix(matEntries,ncol=4)
      [,1] [,2] [,3] [,4]
[1,]     1     5     9    13
[2,]     2     6    10    14
[3,]     3     7    11    15
[4,]     4     8    12    16
> matrix(matEntries,ncol=2)
      [,1] [,2]
[1,]     1     9
[2,]     2    10
[3,]     3    11
[4,]     4    12
[5,]     5    13
[6,]     6    14
[7,]     7    15
[8,]     8    16
> matrix(matEntries,ncol=8)
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8]
[1,]     1     3     5     7     9    11    13    15
[2,]     2     4     6     8    10    12    14    16
```

Note that the entries get rolled up column by column.

Entry-wise product is with `*`, and matrix multiplication is with `%*%`. Matrices can be inverted with `solve`.

Selecting items from arrays/matrices

Suppose that we have an array `arr` consisting of the numbers from 1 through 10.

```
> arr = 1:10
```

We can the 3rd element of the array by doing `arr[3]`, with square brackets:

```
> arr[3]
[1] 3
```

If we want to create a new array with the 1st, 5th, 8th, and 10th entry, then we put a *list* with the indices 1, 5, 8, 10 within the square brackets:

```
> arr[c(1,5,8,10)]
[1] 1 5 8 10
```

(Question: Why did we write `c(1,5,8,10)` and not just `1,5,8,10`? Try without the `c` and see what happens.)

Now suppose that we have a long list of the numbers from 1 through 100, and we want to get from it another array with the entries that are multiples of 7. How do we do this? First, we notice that if we pass an array of *boolean* (true/false) values to the array, we get the entries that correspond to true:

```
> arr2 = c("a","b","c","d","e","f")
> arr2[c(TRUE, FALSE, FALSE, TRUE, TRUE, FALSE)]
[1] "a" "d" "e"
```

Furthermore, if we compare an *array* with a single item, we get a list of entry-by-entry comparison. For example

```
> arr3 = c(-19,27,43,0,2,-2,-8,10)
> arr3 >= 8
[1] FALSE TRUE TRUE FALSE FALSE FALSE TRUE
```

Here we are taking each entry of `arr3` and asking whether it is ≥ 8 or not. Then we get an array with `TRUE` if the element is indeed ≥ 8 , or `FALSE` if not.

Now we can combine these last ideas to extract our array of multiples of 7.

```
> arr4 = 1:100
> whichmultseven = (arr4%%7 == 0)
> arr4[whichmultseven]
[1] 7 14 21 28 35 42 49 56 63 70 77 84 91 98
```

In the first line we construct the array of numbers from 1 through 100.

In the second line, we construct an array of TRUE/FALSE values, where each value tells us if the corresponding entry of `arr4` is a multiple of 7 or not.

In the last line, we pass this array of TRUE/FALSE values to the square brackets in `arr4` to get the ones that are indeed multiples of 7.

We can condense all of this in a single line:

```
arr4[arr4%%7 == 0]
[1] 7 14 21 28 35 42 49 56 63 70 77 84 91 98
```

Now when we have matrices, we must pass two values within square brackets to get an entry back:

```
> mat1 = matrix(arr4, nrow=10)
> mat1
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
[1,]    1   11   21   31   41   51   61   71   81   91
[2,]    2   12   22   32   42   52   62   72   82   92
[3,]    3   13   23   33   43   53   63   73   83   93
[4,]    4   14   24   34   44   54   64   74   84   94
[5,]    5   15   25   35   45   55   65   75   85   95
[6,]    6   16   26   36   46   56   66   76   86   96
[7,]    7   17   27   37   47   57   67   77   87   97
[8,]    8   18   28   38   48   58   68   78   88   98
[9,]    9   19   29   39   49   59   69   79   89   99
[10,]  10   20   30   40   50   60   70   80   90  100
> mat1[8,9]
[1] 88
```

Here we can be more creative and get *submatrices*. For example, suppose that we just want the 2nd and 3rd column and the 5th, 9th and 10th row. Then we can pass those

as *arrays* in the square brackets:

```
> mat1[c(2,3),c(5,9,10)]
      [,1] [,2] [,3]
[1,]   42   82   92
[2,]   43   83   93
```

What if we want a few rows, but *all* columns? Then we leave the column index blank:

```
> mat1[c(5,9), ]
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
[1,]     5   15   25   35   45   55   65   75   85   95
[2,]     9   19   29   39   49   59   69   79   89   99
```

Similarly, if we want just a few columns but *all* rows, we leave the row index blank:

```
> mat1[ , c(1,5,7)]
      [,1] [,2] [,3]
[1,]     1   41   61
[2,]     2   42   62
[3,]     3   43   63
[4,]     4   44   64
[5,]     5   45   65
[6,]     6   46   66
[7,]     7   47   67
[8,]     8   48   68
[9,]     9   49   69
[10,]    10   50   70
```

Reading docstrings and function parameters

As we saw above, we can look for the documentation of different things with either the `help` command or using the searchbar in the Help window. For example, if I type `help(runif)` I get the following (massive) docstring

The Uniform Distribution

Description

These functions provide information about the uniform distribution on the interval from min to max. `dunif` gives the density, `punif` gives the

distribution function qunif gives the quantile function and runif generates random deviates.

Usage

```
dunif(x, min = 0, max = 1, log = FALSE)
```

```
punif(q, min = 0, max = 1, lower.tail = TRUE, log.p = FALSE)
```

```
qunif(p, min = 0, max = 1, lower.tail = TRUE, log.p = FALSE)
```

```
runif(n, min = 0, max = 1)
```

Arguments

x, q vector of quantiles.

p vector of probabilities.

n number of observations. If $\text{length}(n) > 1$, the length is taken to be the number required.

min, max lower and upper limits of the distribution. Must be finite.

log, log.p logical; if TRUE, probabilities p are given as $\log(p)$.

lower.tail logical; if TRUE (default), probabilities are $P[X \leq x]$, otherwise, $P[X > x]$.

Details

If min or max are not specified they assume the default values of 0 and 1 respectively.

The uniform distribution has density

$$f(x) = 1/(\text{max}-\text{min})$$

for $\text{min} \leq x \leq \text{max}$.

For the case of $u := \text{min} == \text{max}$, the limit case of $X == u$ is assumed, although there is no density in that case and dunif will return NaN (the error condition).

runif will not generate either of the extreme values unless $\text{max} = \text{min}$ or $\text{max}-\text{min}$ is small compared to min, and in particular not for the default arguments.

Value

dunif gives the density, punif gives the distribution function, qunif gives the quantile function, and runif generates random deviates.

The length of the result is determined by n for runif, and is the maximum of the lengths of the numerical arguments for the other functions.

The numerical arguments other than n are recycled to the length of the result. Only the first elements of the logical arguments are used.

Note

The characteristics of output from pseudo-random number generators (such as precision and periodicity) vary widely. See `.Random.seed` for more information on R's random number generation algorithms.

References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) The New S Language. Wadsworth & Brooks/Cole.

See Also

RNG about random number generation in R.

Distributions for other standard distributions.

Examples

```
u ← runif(20)
```

```
/## The following relations always hold :
```

```
punif(u) == u
```

```
dunif(u) == 1
```

```
var(runif(10000))  $\sim$  1/12 = .08333
```

Let's focus on `runif(n, min = 0, max = 1)`. If we want to use it to generate 10 uniformly distributed numbers between 0 and 1, we must do

```
> runif(10)
```

Note that the docstring showed that `runif` has 3 parameters `n, min, max`, but in the call above we only used one. This is because `min` and `max` have *default values* of 0 and 1, respectively. We can identify when a parameter has a default value when it is defined like `param=something` in the docstring.

Exercises to try:

Most exercises are just about syntax (so solved by looking at the documentation or by googling). These are a bit more fun:


7. Using the functions `runif()` and `trunc()` simulate 20 die throws.
8. Generate a randomisation list of size 150 for three treatment groups "Placebo", "Drug A" and "Drug B".

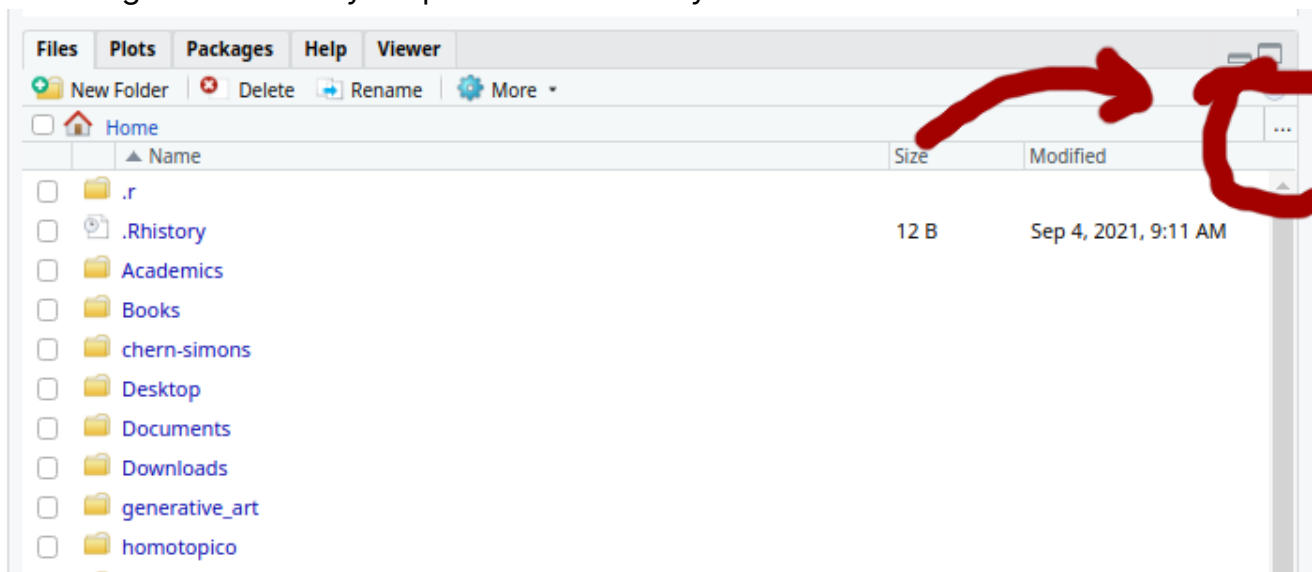
Loading data and dataframes

Before we even try to load data in R and RStudio, we have to let it know what the working directory is. The **working directory** is the place where R will store files and where it will look for data.

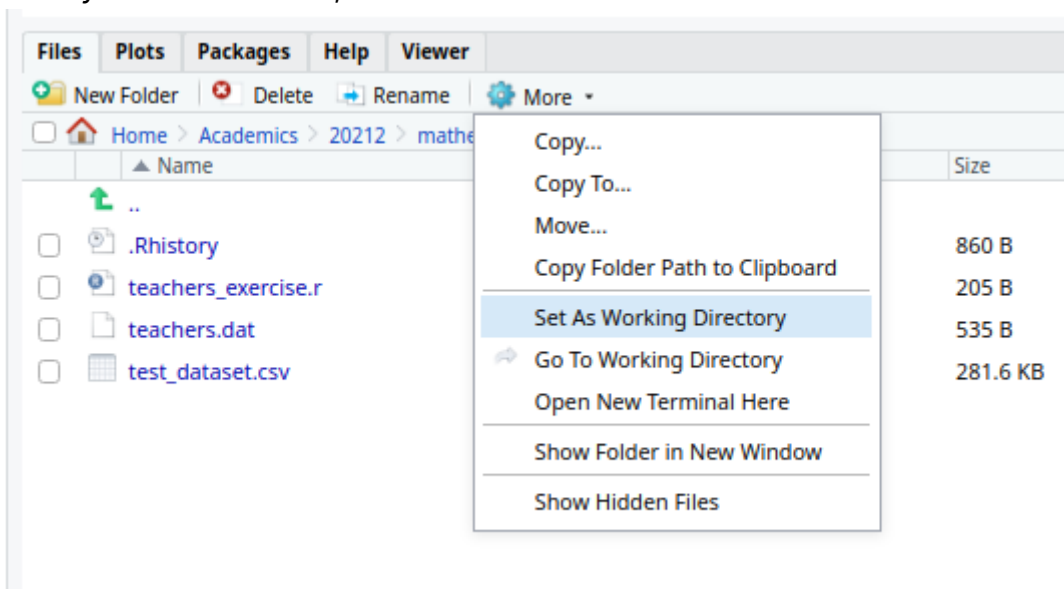
In barebones R (i.e. the command console) we can set the working directory with

```
> setwd("path/to/directory")
```

However, in RStudio we can set it manually. In the lower right pane, click on the **Files** tab. Here you can either navigate directly to the folder you want, or use the  icon on the right and select your preferred directory.



After you've done that, click on the Gear icon and then on **Set as Working Directory**.



Once you do this, you've told R *where* to look for files.

Now download the dataset (either from ItsLearning or from [this link](#)) and *make sure that the dataset is in the working directory*.

Think of a dataframe like an Excel spreadsheet.

We load dataframes from textfiles with `read.table()`, and if the data is in CSV format (that's Comma-Separated Values), we would use `read.csv()`. In this case, we are interested in the dataset `test_dataset.csv`. We load it into a variable called `dataset`.

```
> dataset = read.csv("test_dataset.csv")
```

We can check the first few rows of the dataset with the `head()` function:

```
> head(dataset)
```

	name	pokedex_number	abilities	typing	hp	attack
1	Bulbasaur	1	Overgrow~Chlorophyll	Grass~Poison	45	49
2	Ivysaur	2	Overgrow~Chlorophyll	Grass~Poison	60	62
3	Venusaur	3	Overgrow~Chlorophyll	Grass~Poison	80	82
4	Venusaur Gmax	3	Overgrow~Chlorophyll	Grass~Poison	80	82
5	Venusaur Mega	3	Thick Fat	Grass~Poison	80	100
6	Charmander	4	Blaze~Solar Power	Fire	39	52

Calling the `str()` function also gives us an overview of the dataframe:

```
> str(dataset)
```

```
'data.frame':  1118 obs. of  49 variables:
 $ name          : chr  "Bulbasaur" "Ivysaur" "Venusaur"
 "Venusaur Gmax" ...
 $ pokedex_number : int   1  2  3  3  3  4  5  6  6  6 ...
 $ abilities      : chr  "Overgrow~Chlorophyll"
 "Overgrow~Chlorophyll" "Overgrow~Chlorophyll" ...
 $ typing         : chr  "Grass~Poison" "Grass~Poison"
 "Grass~Poison" "Grass~Poison" ...
 $ hp             : int   45  60  80  80  80  39  58  78  78  78 ...
 $ attack         : int   49  62  82  82  100  52  64  84  84  130 ...
 $ defense        : int   49  63  83  83  123  43  58  78  78  111 ...
 $ special_attack : int   65  80  100  100  122  60  80  109  109  130
 ...
 $ special_defense : int   65  80  100  100  120  50  65  85  85  85 ...
 $ speed          : int   45  60  80  80  80  65  80  100  100  100 ...
 $ height         : int    7  10  20  240  24  6  11  17  280  17 ...
 $ weight         : int   69  130  1000  10000  1555  85  190  905
 10000 1105 ...
 $ genus          : chr  "Seed Pok\303\251mon" "Seed
 Pok\303\251mon" "Seed Pok\303\251mon" "Seed Pok\303\251mon" ...
 ...
```

With this we see that this is a list of Pokémon.^{[1][2]}

We can get a column from `df$columnname`, and we can return (and change) the column names by querying `names()`:

```
> dataset$name
[1] Bulbasaur          Ivysaur
[3] Venusaur           Venusaur Gmax
[5] Venusaur Mega      Charmander
[7] Charmeleon         Charizard
[9] Charizard Gmax     Charizard Mega X
[11] Charizard Mega Y   Squirtle
[13] Wartortle          Blastoise
[15] Blastoise Gmax     Blastoise Mega
[17] Caterpie           Metapod
[19] Butterfree         Butterfree Gmax
[21] Weedle             Kakuna
[23] Beedrill           Beedrill Mega
[25] Pidgey             Pidgeotto
...
> names(dataset)
[1] "name"                "pokedex_number"
[3] "abilities"           "typing"
[5] "hp"                  "attack"
[7] "defense"             "special_attack"
[9] "special_defense"     "speed"
[11] "height"              "weight"
[13] "genus"               "gen_introduced"
[15] "female_rate"         "genderless"
[17] "baby_pokemon"        "legendary"
[19] "mythical"            "is_default"
[21] "forms_switchable"    "base_experience"
[23] "capture_rate"        "egg_groups"
[25] "egg_cycles"          "base_happiness"
[27] "can_evolve"          "evolves_from"
[29] "primary_color"       "shape"
[31] "number_pokemon_with_typing" "normal_attack_effectiveness"
[33] "fire_attack_effectiveness" "water_attack_effectiveness"
[35] "electric_attack_effectiveness" "grass_attack_effectiveness"
[37] "ice_attack_effectiveness" "fighting_attack_effectiveness"
[39] "poison_attack_effectiveness" "ground_attack_effectiveness"
[41] "fly_attack_effectiveness" "psychic_attack_effectiveness"
[43] "bug_attack_effectiveness" "rock_attack_effectiveness"
[45] "ghost_attack_effectiveness" "dragon_attack_effectiveness"
```

```
[47] "dark_attack_effectiveness"      "steel_attack_effectiveness"
[49] "fairy_attack_effectiveness"
```

This dataset has way too many columns (and if we know the rock-paper-scissors rules between Pokémon types, we know the attack effectiveness so columns 32-49 are a lot of superfluous information), so we want to select just a few. Say that we want to keep only the first 20 columns. Thinking of this dataset as a matrix, from the above we can do

```
> dataset = dataset[ , 1:20]
```

i.e. selecting *all* rows and our specified columns.

Now suppose that we only wanted to keep the columns `name`, `hp`, and `typing`. Here we can pass a list of *column names* to the square brackets and it will work:

```
> head(dataset[ , c("name", "hp", "typing")])
      name hp      typing
1  Bulbasaur 45 Grass~Poison
2   Ivysaur 60 Grass~Poison
3   Venusaur 80 Grass~Poison
4 Venusaur Gmax 80 Grass~Poison
5 Venusaur Mega 80 Grass~Poison
6  Charmander 39      Fire
...
```

Note that the order of the columns is the same as the order of the list we passed, even though in the original dataset, `typing` goes before `hp`.

We can add a column to the right with `cbind(df, newCol)`, and a row below with `rbind(df, newRow)`.

-
1. Did you know that Pokémon is the [highest-grossing media franchise](#)? And most of the profits come from merch.↩
 2. I got this dataset from [Kaggle](#). There's many others to play with!↩